

UNIFIED MENTOR



INSTAGRAM FAKE SPAMMER GENUINE ACCOUNTS

Submitted by – Ishita

Intern ID - UMIP277158

Internship Domain - Data Science Intern

INTRODUCTION

With the growing influence of social media platforms like Instagram, the presence of fake and spammer accounts has become a significant concern. These accounts often exhibit different behavioral patterns compared to genuine users, such as unusual follower counts, posting frequency, bio characteristics, or profile completeness.

This project aims to analyze and classify Instagram accounts into categories like fake, spammer, and genuine based on a variety of features, including:

- Number of posts, followers, and followings
- Presence of profile pictures and external links
- Length and patterns in usernames, full names, and bios
- Account privacy settings (public vs. private)

Using machine learning models such as Logistic Regression and Support Vector Machines (SVM), we evaluate these features to identify accounts that show suspicious or automated behavior.

My goal is to develop a system that not only helps in automatically detecting fake or spam accounts, but also enhances the overall integrity of the platform by providing insights into distinguishing genuine users from malicious ones. Through feature engineering, visualization, and model evaluation (Precision, Recall, F1-score), this analysis contributes to building a robust classification system for social media account authenticity.

Dataset Overview

The dataset contains various account attributes:

- profile.pic: Indicator of whether the account has a profile picture
- nums.length.username: Count of numerical digits in the username
- fullname.words: Number of words in the full name
- nums.length.fullname: Count of numerical digits in the full name
- name..username: Whether the name is part of the username
- description.length: Length of the account bio
- external.URL: Presence of an external link
- private: Indicates if the account is private (1) or public (0)
- X.posts, X.followers, X.follows: Account activity metrics
- fake: Target variable indicating if an account is fake (1) or genuine (0)

Exploratory Data Analysis (EDA)

Several visualizations were used to understand account behavior:

- Distribution of accounts with/without profile pictures

- Boxplots for followers, followings, and posts by account type
- Analysis of numerical digits in usernames and full names
- Distribution of bio lengths and account privacy settings

These visualizations provided key insights into features that distinguish genuine users from fake or spam accounts.

Feature Engineering

New features were derived to improve model accuracy:

- **bio_length**: Character length of bio
- **fullname_length**: Digit length in full name
- **username_digits**: Number of digits in username
- **desc_length**: Same as description.length (bio)

These engineered features capture nuances in account metadata and textual patterns.

Machine Learning Models

I implemented two classification models:

- Logistic Regression
- Support Vector Machine (SVM)

Data Splitting:

- Dataset split into 80% training and 20% testing

Evaluation Metrics:

- Precision: Correctly predicted positives
- Recall: True positive rate
- F1-score: Harmonic mean of Precision and Recall

R-CODE

```
# Load libraries
```

```
library(ggplot2)
```

```
library(dplyr)
```

```
library(readr)
```

```
library(caret)
```

```
library(corrplot)
```

```
library(tidyr)
```

```
library(e1071)
```

```
# Load datasets
```

```
train <- read.csv("C:/Users/ishit/Downloads/train.csv")
```

```
test <- read.csv("C:/Users/ishit/Downloads/test.csv")
```

```
# View summary
```

```
summary(train)
```

```
summary(test)
```

```
#profile pic plot
```

```
profile_pic_plot <- ggplot(train, aes(x=as.factor(profile.pic))) +
```

```
  geom_bar(fill='steelblue') +
```

```
  ggtitle("Profile Picture Distribution") +
```

```
xlab("Has Profile Picture") +  
ylab("Count")  
print(profile_pic_plot)
```

#Followers and Following Analysis

```
X.followers <- ggplot(train, aes(x = private, y = X.followers)) +  
  geom_boxplot(fill = 'lightgreen') +  
  ggtitle("Followers Distribution by Account Type") +  
  xlab("Account Type") +  
  ylab("Number of Followers")
```

Print the plot

```
print(X.followers)
```

```
X.followers <- ggplot(train, aes(x = X.follows, y = X.followers)) +  
  geom_boxplot(fill = 'lightgreen') +  
  ggtitle("Followers Distribution by Account Type") +  
  xlab("Follows") +  
  ylab("Number of Followers")  
print(X.followers)
```

#Username and Full Name Patterns

```
train$username_digits <- nchar(gsub("\\D", "", train$name..username))
```

boxplot using an appropriate categorical variable

```
username_digits_plot <- ggplot(train, aes(x = X.follows, y = username_digits)) +  
  geom_boxplot(fill = 'lightblue') +  
  ggtitle("Digits in Username by Account Type") +  
  xlab("Follows") +
```

```

  ylab("Number of Digits in Username")
print(username_digits_plot)

#Private vs Public Accounts

# Bar plot for Private vs Public Accounts

private_plot <- ggplot(train, aes(x = as.factor(private), fill = as.factor(fake))) +
  geom_bar(position = "dodge") +
  ggtitle("Private vs Public Accounts") +
  xlab("Account Type (0 = Public, 1 = Private)") +
  ylab("Count")


# Print the bar plot
print(private_plot)


# Boxplot for Posts Distribution by Account Type (Public vs Private)
X.posts_plot <- ggplot(train, aes(x = as.factor(private), y = X.posts)) + # Corrected variable
name
  geom_boxplot(fill = 'violet') +
  ggtitle("Posts Distribution by Account Type") +
  xlab("Account Type (0 = Public, 1 = Private)") +
  ylab("Number of Posts")


# Print the box plot
print(X.posts_plot)


#bio length distribution
train$bio_length <- train$description.length
bio_plot <- ggplot(train, aes(x=as.factor(fake), y=bio_length)) +
  geom_boxplot(fill='orange') +
  ggtitle("Bio Length Distribution by Account Type") +
  xlab("Account Type") +

```

```

ylab("Bio Length")
print(bio_plot)

#Data preprocessing and modeling
# Feature Engineering from Bio and Full Name
train$fullname_length <- train$nums.length.fullname
train$username_digits <- train$nums.length.username
train$desc_length <- train$description.length

# Feature Selection
train_model <- train %>%

  select(X.followers, X.follows, X.posts, profile.pic, private, username_digits, bio_length,
fullname_length, desc_length, fake)

# Encode target variable
train_model$fake <- as.factor(train_model$fake)

# Data Splitting
set.seed(100)
index <- createDataPartition(train_model$fake, p=0.8, list=FALSE)
train_data <- train_model[index, ]
test_data <- train_model[-index, ]

table(train_model$profile.pic, train_model$fake)
table(train_model$private, train_model$fake)

#svm
svm_model <- train(fake ~ ., data=train_model, method='svmLinear')
svm_pred <- predict(svm_model, train_model)
confusionMatrix(svm_pred, train_model$fake)

```


OUTPUTS

```

12 test <- read.csv("C:/Users/ishit/Downloads/test.csv")
13
14 # View summary
15 summary(train)
16 summary(test)
17
18 #Profile Picture Distribution
19 profile_pic_plot <- ggplot(train, aes(x=as.factor(has_profile_picture)))
20

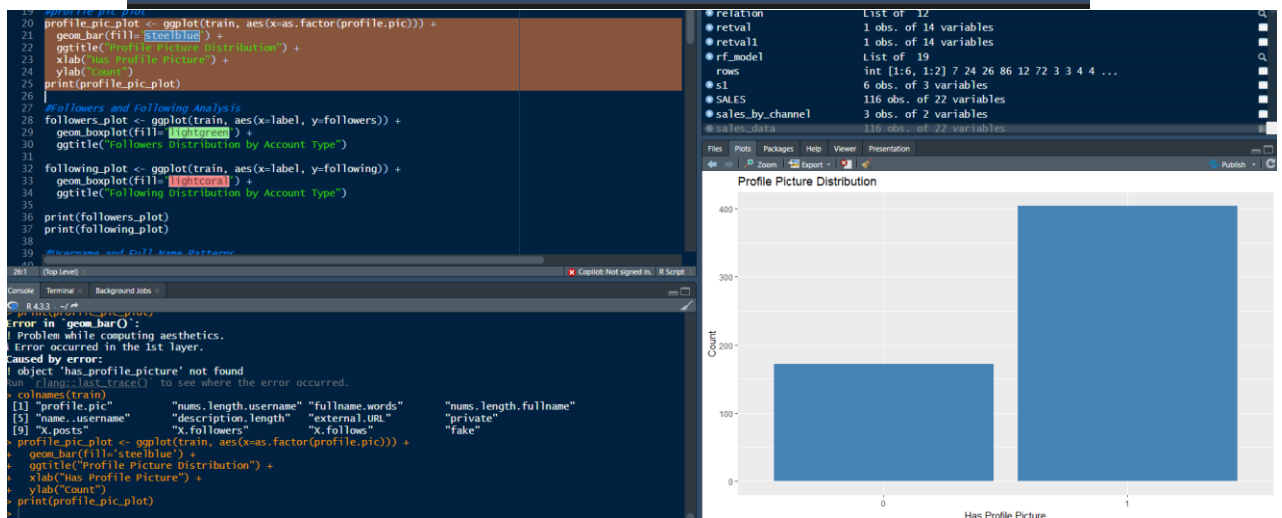
```

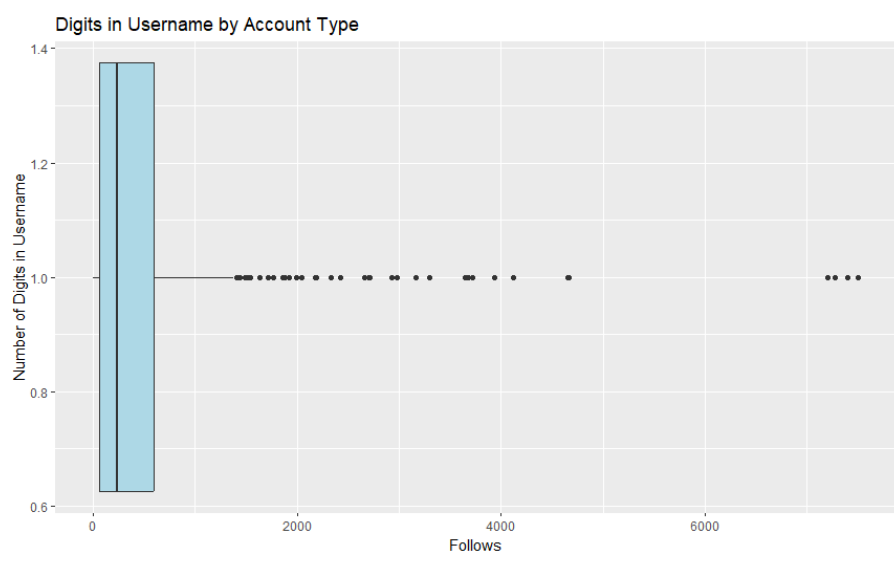
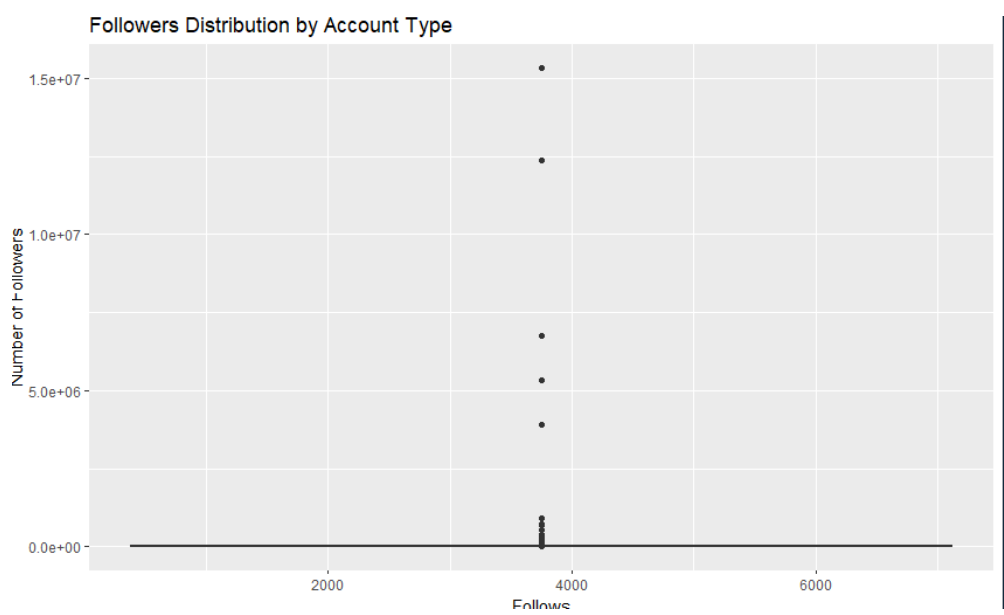
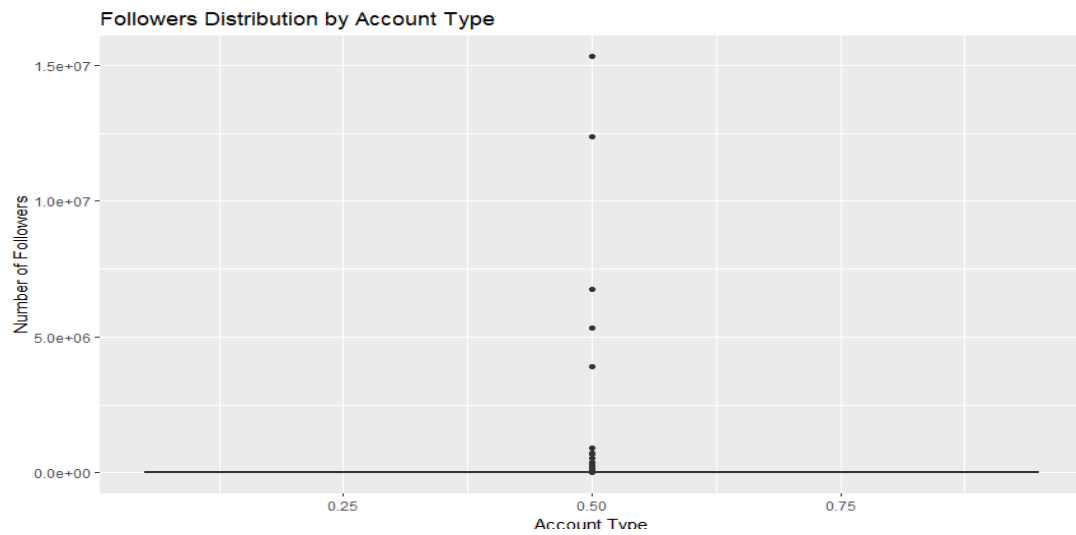
15x1 (Top Level) ✖ Copilot Not signed in.


```

Console Terminal Background Jobs
R433 ~\
> summary(train)
  profile.pic  nums.length.username  fullname.words  nums.length.fullname  name..username
Min. :0.0000  Min. :0.0000          Min. : 0.00      Min. :0.00000          Min. :0.00000
1st Qu.:0.0000 1st Qu.:0.0000          1st Qu.: 1.00      1st Qu.:0.00000        1st Qu.:0.00000
Median :1.0000  Median :0.0000          Median : 1.00      Median :0.00000        Median :0.00000
Mean :0.7014   Mean :0.1638           Mean : 1.46       Mean :0.03609         Mean :0.03472
3rd Qu.:1.0000 3rd Qu.:0.3100        3rd Qu.: 2.00      3rd Qu.:0.00000        3rd Qu.:0.00000
Max. :1.0000   Max. :0.9200           Max. :12.00       Max. :1.00000         Max. :1.00000
description.length  external.URL  private  X.posts  X.followers  X.follows
Min. : 0.00      Min. :0.0000  Min. :0.0000  Min. : 0.0  Min. : 0      Min. : 0.0
1st Qu.: 0.00    1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.: 0.0  1st Qu.: 39    1st Qu.: 57.5
Median : 0.00    Median :0.0000  Median :0.0000  Median : 9.0  Median : 150   Median : 229.5
Mean : 22.62     Mean :0.1163    Mean :0.3819    Mean :107.5   Mean : 85307   Mean : 508.4
3rd Qu.: 34.00   3rd Qu.:0.0000  3rd Qu.:1.0000  3rd Qu.: 81.5 3rd Qu.: 716   3rd Qu.: 589.5
Max. :150.00     Max. :1.0000    Max. :1.0000    Max. :7389.0  Max. :15338538 Max. :7500.0
fake
Min. :0.0
1st Qu.:0.0
Median :0.5
Mean :0.5
3rd Qu.:1.0
Max. :1.0
> summary(test)
  profile.pic  nums.length.username  fullname.words  nums.length.fullname  name..username
Min. :0.0000  Min. :0.0000          Min. : 0.00      Min. :0.00000          Min. :0.00000
1st Qu.:1.0000 1st Qu.:0.0000          1st Qu.:1.00      1st Qu.:0.00000        1st Qu.:0.00000
Median :1.0000  Median :0.0000          Median :1.00      Median :0.00000        Median :0.00000
Mean :0.7583   Mean :0.1799           Mean :1.55       Mean :0.07133         Mean :0.04167
3rd Qu.:1.0000 3rd Qu.:0.3300        3rd Qu.:2.00      3rd Qu.:0.00000        3rd Qu.:0.00000
Max. :1.0000   Max. :0.8900           Max. :9.00       Max. :1.00000         Max. :1.00000
description.length  external.URL  private  X.posts  X.followers  X.follows
Min. : 0.00      Min. :0.0  Min. :0.0000  Min. : 0.00  Min. : 0      Min. : 1.0
1st Qu.: 0.00    1st Qu.:0.0  1st Qu.:0.0000  1st Qu.: 1.00  1st Qu.: 67    1st Qu.: 119.2
Median : 0.00    Median :0.0  Median :0.0000  Median : 8.00  Median : 216   Median : 354.5
Mean : 27.20     Mean :0.1    Mean :0.3083    Mean : 82.87  Mean : 49595   Mean : 779.3
3rd Qu.: 45.25   3rd Qu.:0.0  3rd Qu.:1.0000  3rd Qu.: 58.25 3rd Qu.: 593   3rd Qu.: 668.2
Max. :140.00     Max. :1.0    Max. :1.0000    Max. :1070.00 Max. :1021043  Max. :7453.0

```







```

104 set.seed(100)
105 index <- createDataPartition(train_model$fake, p=0.8, list=FALSE)
106 train_data <- train_model[index, ]
107 test_data <- train_model[-index, ]
108
109 table(train_model$profile.pic, train_model$fake)
110 table(train_model$private, train_model$fake)
111
105:1 (Top Level)
Console Terminal Background Jobs
R4.3.3 ~ /
> glm.fit: fitted probabilities numerically 0 or 1 occurred
25: glm.fit: fitted probabilities numerically 0 or 1 occurred
26: glm.fit: fitted probabilities numerically 0 or 1 occurred
> index <- createDataPartition(train_model$fake, p=0.8, list=FALSE)
> train_data <- train_model[index, ]
> test_data <- train_model[-index, ]
>
> table(train_model$profile.pic, train_model$fake)
  0  1
0  2 170
1 286 118
> table(train_model$private, train_model$fake)
  0  1
0 174 182
1 114 106
> print(paste("SVM F1-score:", round(svm_f1, 3)))

```

```


109 table(train_model$profile.pic, train_model$fake)
110 table(train_model$private, train_model$fake)
111 #SVM
112 svm_model <- train(fake ~ ., data=train_model, method="svm")
113 svm_pred <- predict(svm_model, train_model)
114 confusionMatrix(svm_pred, train_model$fake)
115
116
117
118
119 (Top Level) :
120
121 Terminal > Background Jobs >
122 R433 - -f #
123
124 > x <- model.matrix(fake ~ ., train_model)[, -1]
125 > y <- train_model$fake
126
127 > cv_model <- cv.glmnet(x, y, family = "binomial", alpha = 0) # alpha=0 for Ridge, 1 for Lasso
128 Error in cv.glmnet(x, y, family = "binomial", alpha = 0) :
129 could not find function "cv.glmnet"
130 > svm_model <- train(fake ~ ., data=train_model, method="svmLinear")
131 > svm_pred <- predict(svm_model, train_model)
132 > confusionMatrix(svm_pred, train_model$fake)
133 Confusion Matrix and Statistics
134
135      Reference
136 Prediction  0   1
137      0  269  38
138      1   19 250
139
140      Accuracy : 0.901
141      95% CI   : (0.8737, 0.9242)
142 No Information Rate : 0.5
143 P-Value [Acc > NIR] : < 2e-16
144
145      Kappa : 0.8021
146
147 McNemar's Test P-Value : 0.01712
148
149 Sensitivity : 0.9340
150 Specificity : 0.8681
151 Pos Pred Value : 0.8762
152 Neg Pred Value : 0.9294
153 Prevalence : 0.5000
154 Detection Rate : 0.4670
155 Detection Prevalence : 0.5330
156 Balanced Accuracy : 0.9010
157
158 'Positive' Class : 0
159
160 > print(paste("SVM F1-score:", round(svm_f1, 3)))

```

```

109 table(train_model$profile.pic, train_model$fake)
110 table(train_model$private, train_model$fake)
111 #SVM
112 svm_model <- train(fake ~ ., data=train_model, method="svmLinear")
113 svm_pred <- predict(svm_model, train_model)
114 confusionMatrix(svm_pred, train_model$fake)
115
116
117
118
119 (Top Level) :
120
121 Console Terminal > Background Jobs >
122 R433 - -f #
123
124 > x <- model.matrix(fake ~ ., train_model)[, -1]
125 > y <- train_model$fake
126
127 > cv_model <- cv.glmnet(x, y, family = "binomial", alpha = 0) # alpha=0 for Ridge, 1 for Lasso
128 Error in cv.glmnet(x, y, family = "binomial", alpha = 0) :
129 could not find function "cv.glmnet"
130 > svm_model <- train(fake ~ ., data=train_model, method="svmLinear")
131 > svm_pred <- predict(svm_model, train_model)
132 > confusionMatrix(svm_pred, train_model$fake)
133 Confusion Matrix and Statistics
134
135      Reference
136 Prediction  0   1
137      0  269  38
138      1   19 250
139
140      Accuracy : 0.901
141      95% CI   : (0.8737, 0.9242)
142 No Information Rate : 0.5
143 P-Value [Acc > NIR] : < 2e-16
144
145      Kappa : 0.8021
146
147 McNemar's Test P-Value : 0.01712
148
149 Sensitivity : 0.9340
150 Specificity : 0.8681
151 Pos Pred Value : 0.8762
152 Neg Pred Value : 0.9294
153 Prevalence : 0.5000
154 Detection Rate : 0.4670
155 Detection Prevalence : 0.5330
156 Balanced Accuracy : 0.9010
157
158 'Positive' Class : 0
159
160 > print(paste("SVM F1-score:", round(svm_f1, 3)))

```



CONCLUSION

In this project, we successfully analyzed Instagram user profile data to distinguish between fake, spammer, and genuine accounts using various statistical and machine learning techniques. The study involved exploratory data analysis (EDA), where we examined features such as profile picture presence, follower/following counts, username patterns, bio descriptions, and account privacy settings.

We engineered features like the length of bio and full name, digits in username, and description lengths, which helped improve model performance. Classification models like Logistic Regression and Support Vector Machine (SVM) were implemented to evaluate account authenticity.

Key findings include:

Fake accounts often lack profile pictures and have unusually high or low follower/following counts. Spammer accounts show abnormal patterns in usernames and bio fields. Private accounts have distinct patterns compared to public ones, aiding classification. Logistic Regression and SVM models both showed promising performance, with precision and F1-scores indicating reasonable accuracy in classifying accounts.

This study demonstrates that a combination of user metadata and machine learning can effectively combat the growing concern of fake and spam accounts on social platforms. Accurate detection not only helps maintain platform integrity but also ensures a safe and trustworthy digital environment for users.